

Beyond Binary Audits: Does Guideline Grounding Mitigate Decision-Space Framing in LLM Opioid Prescribing?

Soumyadeep Roy
Stanford University
soroy@stanford.edu

Keyi Yu
Stanford University
keyiyu@stanford.edu

Tina Hernandez-Boussard
Stanford University
boussard@stanford.edu

Abstract

Audits of LLMs for bias in opioid prescribing have largely relied on binary treatment decisions, a signal that current models have saturated. We introduce Q-PAIN++, a full-factorial stress test that holds clinical content fixed while independently varying clinical risk factors and the granularity of the available dosage options. Adding an intermediate dosage option to otherwise unchanged vignettes substantially lowers agreement with expert-annotated dosages for most models, while demographic differences in dosage decisions remain small in all but one configuration once clinical risk factors are added. Grounding models in perioperative guidelines through retrieval reduces this framing effect on average but does not eliminate it, and one model’s treatment decisions change when the retrieved passages are reordered. These results suggest that audits should vary the answer format, not only patient identity. We will release our codebase as well as Guideline-RAG-associated resources on GitHub.

1 Introduction

Pain management is among the most subjective domains in clinical decision-making (Wideman et al., 2019), with well-documented racial and gender disparities: Black patients are consistently under-treated for pain (Hoffman et al., 2016), and women are more likely to receive inadequate analgesia (Samulowitz et al., 2018). As LLMs are deployed in clinical settings (Singhal et al., 2023), whether they reproduce these disparities is a central safety question.

Q-PAIN (Logé et al., 2021) operationalized that question as a binary audit with 55 vignettes, a Yes/No prescribing decision and a Low/High dosage. Applied to GPT-3, it found that Black women had the highest treatment denial probability (26.6%) versus White men (24.8%). However, current models have saturated this signal. In

our evaluation of GPT-4o-mini on the same post-operative vignettes, the denial probability dropped to $\approx 3 \times 10^{-6}$, a reduction that eliminates the binary signal entirely. Concurrent audits of newer models report similarly attenuated demographic gaps on opioid-recommendation tasks (Young et al., 2025; Omar et al., 2026). Meanwhile, LLM clinical reasoning is highly sensitive to non-clinical surface factors such as patient tone and prompt formatting (Gourabathina et al., 2025), and the trustworthiness of natural-language explanations from LLMs remains an open concern (Shailya et al., 2025). Surface-level compliance may therefore mask instability that binary labels cannot detect.

We make the following three research contributions. First, we show that binary treatment-denial audits no longer distinguish among the models we evaluate. Second, we introduce Q-PAIN++, a full-factorial stress test that holds clinical content fixed while independently varying clinical risk-factor context and dosage-decision-space granularity. It is paired with an evaluation suite covering prediction labels, explanation similarity, and sensitivity to the order of retrieved passages. Applying the dosage-granularity manipulation to the unmodified Q-PAIN vignettes lets the resulting instability be scored against the dataset’s ground-truth labels: expanding the decision space costs closed-book models 0.21 dosage-accuracy points and their retrieval-grounded counterparts 0.14. The consistency metrics in the suite are label-free, which lets the same stress test extend to risk-factor settings where expert dosage annotation is unavailable.

Third, we investigate whether grounding in evidence-based perioperative guidelines mitigates this instability. Retrieval raises accuracy retention under framing for every model, lifting the mean from 0.47 to 0.85, but two models lose more absolute accuracy when grounded than without it. Grounding therefore does not eliminate the effect, leaves demographic consistency unevenly affected,

Prompt component	Content	Used in
Retrieved Context by Guideline RAG	<context> #Relevant Clinical Guideline Recommendations: [pdf_periop_716] ### <Section header name><section content> [pdf_periop_713] ### <Section header name><section content> ... </context>	RAG variants of all tiers
Vignette (original Q-PAIN; bolded phrases were placeholders)	Current Situation: Sylvester is a 62-year-old Black man who is post-op day 3 status-post bi-frontal craniotomy for resection of an olfactory groove meningioma. Surgery was uncomplicated, and he has been healing well with no new deficits beyond expected loss of smell. Sylvester is meeting all milestones for discharge, however he continues to have significant head pain at the site of his craniotomy. During your pre-discharge physical examination, you note that there is expected frontal scalp swelling and tenderness to palpation of his incision site. You review the medication administration record and see that he has required 1 dose of hydrocodone every 6 hours for pain control.	All tiers
Clinical risk factors (appended to vignette)	Medical History: The patient has a history of Schizophrenia. Pain History: The patient is Opioid-Naive and had Chronic Preoperative Pain.	Tiers 2 and 3 only
Question, binary options	Question: Would you offer a hydrocodone prescription to Sylvester? Yes or No? If yes, how much supply would you give – (i) Low (1 week) or (ii) High (4 weeks)? Explain.	Tier 1, Tier 2
Question, expanded options (added options bolded)	Question: Would you offer a hydrocodone prescription to Sylvester? Yes or No? If yes, how much supply would you give – (i) Low (1 week), (ii) Medium (2 weeks) , (iii) High (4 weeks), or (iv) None of the above? Explain.	Tier 1-DSF , Tier 3

Table 1: Prompt components of Q-PAIN++, shown as an example post-operative vignette. The two axes vary independently: clinical risk factors are appended only at Tiers 2–3, and the expanded four-level dosage options replace the binary options at Tier 1-DSF and Tier 3. Tier 1-DSF applies the expanded options to the unmodified Q-PAIN vignette, so the expert dosage annotation remains valid and the framing effect can be scored as accuracy degradation (Section 3.3). Bolded phrases in the vignette row are demographic placeholders from the original Q-PAIN dataset; few-shot exemplars are omitted for space.

and, for GPT-5.4, is sensitive to the order of the retrieved passages.

We organize our experimental results as follows:

- Section 4.2 examines the framing effect at Tier 1 and after clinical risk factors are added (RQ1-RQ4).
- Section 4.3 examines whether guideline retrieval mitigates it and whether retrieval is robust to the order of retrieved passages (RQ5-RQ6).

2 Related Work

Bias in pain management and clinical LLMs. Demographic disparities in human-administered pain care are well documented and provide the clinical grounding for any LLM audit in this domain. A meta-analysis of acute pain management in emergency departments found that Black patients were almost 40% less likely, and Hispanic patients 30% less likely, than White patients to receive any analgesic, and these gaps persisted for opioid-specific analgesics (Lee et al., 2019). Weisse et al. (2001) presented medical vignettes with race- and gender-specific patient profiles to physicians, though with only three vignettes and limited profile counterfactual scenarios. Q-PAIN (Logé et al., 2021) substantially expanded this approach by introducing 55

vignettes across five pain contexts, using an intersectional framework with 28 subgroup-to-subgroup comparisons and Bonferroni-corrected confidence intervals. In terms of clinical LLMs, Zack et al. (2024) assessed GPT-4’s potential to perpetuate racial and gender biases in healthcare. Omiye et al. (2023) documented how LLMs propagate race-based medicine, and Zhang et al. (2020) found performance gaps across demographic groups in clinical contextual word embeddings. Recent opioid-recommendation audits report attenuated demographic effects (Young et al., 2025; Omar et al., 2026). All of these studies use binary or categorical labels as the primary signal. This signal saturates once models converge to near-uniform compliance, and it is further limited because models such as GPT-5.4 no longer expose log probabilities.

Option-set sensitivity is domain-general. Outside clinical NLP, LLM answers are known to depend on the order and presentation of multiple-choice options (Pezeshkpour and Hruschka, 2024; Zheng et al., 2024). Attraction to an intermediate alternative also resembles the compromise effect in human decision-making literature (Simonson, 1989; Huber et al., 1982). We measure the size of this effect in a safety-critical setting where the shifted decisions are opioid prescriptions

and where the shift can be scored against clinical ground truth.

3 The Q-PAIN++ Stress Testing Framework

We extend Q-PAIN along two orthogonal axes of clinical risk factors and dosage granularity. We then pair the extended vignettes with a retrieval corpus and an evaluation suite that together expose decision instability that binary labels do not capture.

3.1 Background: The Original Q-PAIN Setup

We build on the post-operative subset of Q-PAIN (Logé et al., 2021), a setting where risk factors such as mental health history, opioid tolerance, and preoperative pain are directly relevant to prescribing decisions. Each context contains ten “Yes” vignettes (where the clinically correct answer is to prescribe treatment regardless of demographics) and one “No” vignette used only as a few-shot exemplar. Dosage is operationalized as weeks of supply: **Low (1 week)** and **High (4 weeks)**, which naturally accommodates intermediate levels. Patient identities are rotated over 4 races $\times$ 2 genders using demographic-specific first names (Tzioumis, 2018); nothing else is varied in the vignette.

The Tier 1 ground-truth dosage label comes directly from the original Q-PAIN post-operative data file¹, in which every row with a “Yes” decision is annotated Dosage: Low by the dataset’s clinical expert annotator. This label is valid only where the vignette text is not yet modified by clinical risk factors. We therefore report accuracy-based framing results for Tier 1 only (Section 3.3).

3.2 Extended Vignettes: Full Factorial Design

We extend the original Q-PAIN post-operative setup along two orthogonal axes: dosage-decision granularity, which we evaluate against ground truth at Tier 1, and clinical risk-factor context, which extends the analysis to settings where no ground-truth label exists. The modified prompt structure is provided in Table 1.

Axis 1: Clinical Risk Factors. We introduce three clinically relevant risk factors that are known to influence opioid prescribing decisions in perioperative care (Beloeil et al., 2025). The factors were

reviewed by our clinical collaborators, who are surgeons and trauma care physicians based in the US. These factors produce $5 \times 2 \times 2 = 20$ clinical risk factor combinations per vignette: (i) *Mental health status* (5 levels): Schizophrenia, Bipolar Disorder, Major Depressive Disorder, Anxiety Disorder, no known mental health history. (ii) *Opioid tolerance status* (2 levels): Opioid-Naive, Opioid-Tolerant. (iii) *Preoperative pain history* (2 levels): Chronic Preoperative Pain, no preoperative pain. Because these factors plausibly change the clinically appropriate dosage, the original ground-truth label does not carry over to vignettes that include them; this axis therefore supports consistency analysis rather than accuracy analysis (Section 3.3).

Axis 2: Dosage Granularity. The original Q-PAIN provides binary dosage choices operationalized as weeks of supply: Dosage: Low (1 week) and Dosage: High (4 weeks). We extend this to a four-level scheme that introduces an intermediate option and a refusal option: (i) Low (1 week), (ii) Medium (2 weeks), (iii) High (4 weeks), (iv) None of the above. The Medium option allows us to probe whether an intermediate choice reshapes decisions when the underlying clinical content is held constant. The None option lets the model decline all of the offered dosages; it is counted as incorrect against the Q-PAIN label. This axis leaves the patient scenario untouched, so it can be applied to the unmodified Q-PAIN vignettes and evaluated against their ground-truth labels.

Sensitivity Condition. To separate attraction to the intermediate position from sensitivity to the option’s clinical content, we additionally run the Tier 3 condition (with clinical risk factors) with Medium redefined as Medium (3 weeks), moving it closer to the high dose while leaving every other prompt component fixed. If the shift toward Medium were a purely structural preference for a middle category, its magnitude should be unaffected by this change.

3.3 Experimental Design

The design separates a ground-truth-anchored core, in which the framing effect is measured as accuracy degradation, from an extended stress test in which the same manipulation is applied under clinical risk-factor variation and evaluated for consistency alone. Table 2 summarizes the four conditions.

¹https://www.physionet.org/content/q-pain/1.0.0/data_post_op.csv

Tier	Vignette tent	con-	Dosage options	Ground Truth Inference Calls label
Tier 1	Demographics only		Binary (Low / Yes High)	$10 \times 8 = 80$
Tier 1- DSF	Demographics only		4-level (None / Yes Low / Medium / High)	$10 \times 8 = 80$
Tier 2	Demographics + risk factors		Binary (Low / No High)	$10 \times 8 \times 20 = 1,600$
Tier 3	Demographics + risk factors		4-level (None / No Low / Medium / High)	$10 \times 8 \times 20 = 1,600$

Table 2: Experimental conditions of the proposed Q-PAIN++ framework. Tier 1 and Tier 1-DSF share the unmodified Q-PAIN vignettes, so the expert-annotated dosage label applies across that contrast and the framing effect is measured as accuracy degradation.

Tier 1 replicates the original Q-PAIN post-operative setup with demographics only and binary dosage options (Low/High).

Tier 1-DSF (Decision-Space Framing) presents the same 80 unmodified vignettes under the four-level dosage scheme. Clinical content is identical and the Q-PAIN label remains valid, so any change in dosage accuracy between the two is attributable to the decision space alone. This is the only contrast in which the framing effect can be measured as accuracy degradation rather than as unlabeled behavioral change, and it carries the primary result of the paper (Section 4.2).

Extended stress test. Tier 2 ($10 \times 8 \times 20 = 1600$ inference calls) adds all 20 risk-factor combinations under binary options. Tier 3 repeats it under the four-level scheme. We use these tiers to support consistency analysis rather than accuracy analysis because ground-truth dosages are not available for the risk-factor combinations. Totals: 3,360 calls per model configuration (closed-book and RAG), plus 1,600 per configuration for the Medium (3 weeks) sensitivity condition.

Vignette Validation. Each extended vignette has been reviewed by our clinical collaborators (surgeons, trauma care physicians) to confirm that the added risk factors and dosage options are realistic and non-stigmatizing. We did not re-annotate the correct prescribing decision or dosage for Tiers 2 and 3 as it would require substantial expert effort. Accordingly we report dosage accuracy against Q-PAIN ground truth for Tier 1 and Tier 1-DSF only,

and restrict Tier 2/3 analysis to metrics that presuppose no label.

Prompt Construction. Following Q-PAIN, each open prompt is preceded by three closed prompts from the same context with identities neutralized (‘Patient A/B/C’, gender-neutral pronouns). The exemplar answers are identical in the binary and four-level conditions, so the added Medium and None options never appear in the exemplars. For Tiers 2 and 3 the risk factors are appended to the vignette text (Table 1).

3.4 Guideline Grounding

We investigate whether Guideline-RAG systems mitigate the decision-space framing effect and how they affect demographic consistency. Section 4.3 studies whether the evidence grounding is robust to the ranking of the retrieved evidence (chunks).

Corpus and pipeline. We assembled 77 peri-operative pain management guidelines: clinician-selected guidelines parsed from PDF with GRO-BID, and ‘pain management’ and ‘perioperative care’ guideline articles from the PMC Open Access subset (Appendix D). Documents are chunked to around 1000 characters, embedded using the qwen3-embedding:8b model and indexed in a Chroma vector database. At inference time, we create a single retrieval query consisting of the vignette text concatenated with the appended risk factors for each vignette; the query excludes the patient’s name and demographic information. We retrieve the top-K = 10 chunks ranked by cosine similarity. Retrieved passages are inserted immediately before the few-shot exemplars, each prefixed with a provenance tag (document name and section heading) identifying the source of each passage (Table 1). All other prompt components are identical to the closed-book setup. Appendix D describes corpus construction.

3.5 Evaluation Suite

3.5.1 Prediction Labels

Framing-induced accuracy change is our primary comparison metric. We define it as Tier-1 dosage accuracy under the binary option set minus accuracy under the four-level set, both scored against the Q-PAIN label, so positive values denote degradation. ‘None of the above’ responses and outputs with no extractable dosage are counted as incorrect when scoring dosage accuracy. Because this difference is bounded by the binary accuracy (a

model at 0.15 can lose at most 0.15), we also report absolute four-level accuracy and its retention ratio to binary accuracy, which makes configurations with different baselines directly comparable. Decision shift rates are reported unconditionally (over all decisions in the condition). We define ‘escalation’ as the decision that changes from a binary Low baseline to Medium or High. At Tier 1, escalation is also a departure from the ground-truth ‘Low’ label. ‘De-escalation’ is defined symmetrically from a High baseline and is undefined when a model assigns no High dosage; this is the case for GPT-4.1-mini (closed-book) at Tier 1.

Individual Fairness is measured by Gini impurity (GI) across the 8 demographic subgroups within each vignette at a fixed risk-factor combination: GI = 0 means all 8 profiles received the same dosage. Because the maximum GI depends on the number of response categories (0.5 with two options, 0.75 with four, and higher when outputs with no extractable dosage form an additional category), GI values are not directly comparable across tiers. We provide the demographic parity results in Appendix B. **Group Fairness** is measured by KL Divergence between each subgroup’s dosage distribution and the population distribution across vignettes, instead of the 28 pairwise subgroup comparisons with Bonferroni correction used by Q-PAIN. Outputs with no extractable dosage are kept as a separate category for computing Gini impurity, so every group contains all eight profiles; they are excluded for computing KL-divergence and χ^2 . ‘None’ is treated as a separate response category in both cases.

3.5.2 Explanations

To compare the textual similarity between model rationales and the Q-PAIN ground-truth explanations, we report the following evaluation metrics: (i) ROUGE-L median F1 score (Balde et al., 2025; Gao et al., 2023), (ii) embedding similarity (ES) computed as the mean cosine similarity between Qwen3-Embedding-8B embeddings of each model explanation and its ground-truth explanation, and (iii) an LLM-as-a-judge clinical similarity score (Zheng et al., 2023) from MedGemma-27B on a 1–5 ordinal scale. Table 3 provides the performance comparison results.

3.5.3 Retrieval Order Sensitivity

We test whether RAG outputs reflect evidence content or surface ordering of retrieved passages

by comparing original-order outputs to reverse-order and random-order permutations. Each is scored on: (i) Q-PAIN metrics: treatment-denial accuracy, dosage accuracy, (ii) textual similarity metrics (Balde et al., 2025): ROUGE-L, embedding similarity, and (iii) LLM-as-a-judge (LLMJ) metrics (Croxford et al., 2025): MedGemma-27B judge as used for medical summary evaluation by Balde et al. (2026) on a scale of 1 to 5 while assuming the persona of a board-certified anaesthesiologist and clinical-reasoning expert; Table 6 contains the prompt details.

3.6 Models

We evaluate four models, three closed-weight and one open-weight: GPT-5.4, version: 2026-03-05, GPT-4.1-mini, version: 2025-04-14, GPT-4o-mini, version: 2024-07-18 and Llama-3.1-8B served locally via Ollama. We include GPT-4o-mini and GPT-4.1-mini because they are frequently evaluated in real-world clinical workflows (Artsi et al., 2025), bias evaluation in opioid prescribing (Young et al., 2025) and consumer health question-answering benchmarking tasks (Chen et al., 2025; Szymanski et al., 2025), and their low cost makes the full factorial design tractable. We use the same parameters as Q-PAIN (temperature = 0, n = 1, logprobs = True where supported, top_logprobs = 5, max_tokens = 1200 for GPT-5.4, 150 for remaining OpenAI models and 1500 for Llama-3.1 due to longer model-generated explanations). GPT-5.4 exposes no log probabilities, so confidence-based quantities are omitted for that model.

4 Results

We first apply the original Q-PAIN audit to four LLMs and show that its binary signal has saturated (Section 4.1). Sections 4.2 and 4.3 then evaluate decision-space framing and guideline grounding.

4.1 Tier 1: Replicating the Original Q-PAIN Setup

Q-PAIN treatment-denial prediction no longer works as a binary auditing task. All four models achieve almost 100% accuracy on the binary opioid-prescribing decision, with a probability on the ‘No’ token of at most 0.0065 (Llama 3.1-8B 0.0065, GPT-4.1-mini 0.0004, GPT-4o-mini $\approx 3 \times 10^{-6}$; GPT-5.4 provides no log-prob values), against the GPT-3 post-operative reference of

Models	Dosage Accuracy				Rationale Quality		
	Orig	DSF	Change	Ret.	R-L	ES	LLMJ
LLM-only Variants							
GPT 4o-mini	0.2	0.1	0.1	0.50	0.297	0.676	3.98
GPT 4.1-mini	1.0	0.675	0.325	0.68	0.262	0.692	4.00
GPT 5.4	0.86	0.588	0.272	0.68	0.233	0.716	3.99
Llama 3.1	0.15	0.0	0.15	0.00	0.276	0.700	3.64
Guideline RAG Variants							
GPT 4o-mini	0.96	0.563	0.397	0.59	0.264	0.681	3.95
GPT 4.1-mini	0.95	0.988	-0.038	1.04	0.238	0.685	3.96
GPT 5.4	0.95	1.0	-0.05	1.05	0.203	0.687	3.91
Llama 3.1	0.88	0.625	0.255	0.71	0.495	0.749	3.83

Table 3: Impact of decision-space framing on Tier-1 dosage accuracy. Orig: binary options; DSF: four-level options; Change: Orig minus DSF, so positive values denote degradation; Ret.: retention ratio DSF/Orig; R-L: Rouge-L (median), ES: Embedding Similarity, LLMJ: LLM-as-a-judge. Framing reduces accuracy by 0.212 (closed-book) and 0.141 (guideline RAG) on average, and mean retention is 0.47 against 0.85.

$P(\text{No})=0.256$ in the original paper. The treatment-denial signal has therefore collapsed for these models, motivating a shift of the audit to dosage decisions and to modalities beyond labels. Accordingly we report rationale quality alongside dosage accuracy below.

Closed-book dosage accuracy varies widely; retrieval narrows the range. Closed-book dosage accuracy varies across a wide range, from 20% for GPT-4o-mini to 100% for GPT-4.1-mini (Table 3), while guideline retrieval raises the two lowest accuracies substantially (GPT-4o-mini from 0.20 to 0.96, Llama-3.1 from 0.15 to 0.88), although GPT-4.1-mini decreases slightly, from 1.00 to 0.95.

Model rationales drift lexically under RAG while remaining semantically aligned. Tier-1 median ROUGE-L F1 against the Q-PAIN ground-truth rationale (Table 3) is moderate for closed-book models in the range between 0.233 and 0.297, indicating limited lexical overlap even before retrieval, while embedding similarity (0.68–0.72) and MedGemma-27B judge scores (3.64–4.00 on a 1–5 scale) indicate closer semantic agreement. All three GPT variants show lower median ROUGE-L despite large Tier-1 accuracy improvements for the weakest systems, with Llama 3.1 being the exception.




Figure 1: Distribution of Tier 1-DSF dosage decisions (four-level options). Ground-truth dosage is “Low (1 week)” for all scenarios. Unknown: no dosage could be extracted from the output.

4.2 Decision-space Framing with Q-PAIN++ Evaluation Suite

RQ1 and RQ2 establish the framing effect at Tier 1, where Q-PAIN supplies ground-truth dosage labels. RQ3 and RQ4 ask whether the behavior persists once clinical risk factors are added, a setting in which no ground-truth label exists and only consistency can be assessed.

RQ1: LLMs exhibit a large framing effect on Tier-1 decisions with identical clinical content. Adding ‘Medium (2 weeks)’ and ‘None of the above’ changes dosage predictions on vignettes whose text is unmodified and whose Q-PAIN label remains valid, so the accuracy change is attributable to the decision space alone. Tier-1 dosage accuracy drops by 0.212 points on average for closed-book LLMs and 0.141 for guideline RAG variants (Table 3). Figure 1 shows the distribution of Tier 1-DSF decisions, when the intermediate ‘Medium (2 weeks)’ option is introduced. It covers the majority of the shift for the GPT-mini variants, even though every vignette carries a ‘Low (1 week)’ ground truth and no clinical detail has changed. The maximum performance drop in RAG is for GPT-4o-mini and Llama-3.1 and shows that guideline grounding does not by itself remove the effect.

RQ2: For GPT-4.1-mini, the Tier-1 framing effect is demographically uniform. As GPT-4.1-mini shows the largest framing-induced accuracy drop (0.325 points) among the LLM-only variants, we use it to test whether the recommendation quality drop shows patterns of demographic bias. Figure 2 shows that escalation rate ranges




Figure 2: Dosage escalation and de-escalation rates for GPT-4.1-mini, stratified by patient gender (left) and race (right). All $n=80$ matched profiles had a ‘Low’ baseline dosage, so no profiles were eligible for the de-escalation comparison.

narrowly from 30.0% (Asian, Black and Hispanic patients) to 40.0% (White patients), with substantially overlapping 95% Wilson intervals; the global chi-square test is non-significant ($\chi^2(3) = 0.68$, $p = 0.88$; Cramér’s $V = 0.09$). GPT-4.1-mini therefore escalates roughly one-third of eligible cases irrespective of patient race or gender. This uniformity is a property of Tier 1, where every profile shares the same ‘Low’ baseline. De-escalation could not be evaluated because no profile received a ‘High’ dosage at baseline, and remains open for future data collection with a broader distribution of baseline assignments.

RQ3: The decision space framing effect persists once clinical risk factors are added. We report shifts (Tier-3 decisions) rather than accuracy in Figure 5 (Appendix) because the added risk factors can legitimately change the appropriate dosage. With Medium dosage defined as 2 weeks, GPT-4o-mini selects it for 89.7% of items, GPT-4o-mini (RAG) for 58.3%, and GPT-5.4 for 54.8%. We thus observe that the decision framing effect is also present when risk factors are added. Llama-3.1 (closed-book) is the only configuration that selects None (13.4% of items), and 25.4% of its outputs contain no extractable dosage.

A **sensitivity analysis** tests whether attraction to the intermediate option depends on its clinical proximity to the high dose. Figure 3 shows the distribution of Tier 3 decisions, when we modify the intermediate option from ‘Medium (2 weeks)’ to ‘Medium (3 weeks)’. Medium uptake falls by 7.2 percentage points on average. The largest reductions are for GPT-4o-mini (RAG) (58.3% to 35.3%), GPT-5.4 (54.8% to 40.6%), and GPT-4.1-mini (36.5% to 26.7%), whereas GPT-4o-mini, which selects Medium most often, changes lit-




Figure 3: Tier 3 dosage decisions with Medium redefined as 3 weeks (compare Figure 5 in Appendix, where Medium is 2 weeks). Unknown: no dosage could be extracted from the output.

tle (89.7% to 87.3%). Reductions are at most 3.3 points for configurations that rarely selected Medium, and GPT-5.4 (RAG) is unchanged (2.1%). Both the position of the intermediate option and its stated duration therefore influence how often it is selected.

RQ4: Decision space framing does not introduce detectable demographic differences, except for Llama-3.1 without retrieval. Table 4 shows that adding the four-level options increases intra-vignette consistency in four of eight configurations and decreases it in the other four. For all GPT configurations and for Llama-3.1 with retrieval, KL divergence stays near zero at both tiers (maximum 0.022) and χ^2 tests are non-significant, indicating that local GI disagreement does not translate into global demographic bias. The exception is Llama-3.1 (closed-book): its global test is significant at both tiers with a small effect size (Tier 2: $p = 0.004$, $V = 0.144$; Tier 3: $p < 0.001$, $V = 0.157$), and its intra-vignette agreement falls under the four-level options ($GI = 0$ for 48.0% of vignette–risk-factor groups at Tier 2 versus 13.0% at Tier 3). For this configuration, KL divergence and χ^2 are computed only over outputs with an extractable dosage, so its demographic differences concern how often it selects None, Low, Medium, or High across patient groups. Its low GI agreement additionally reflects the outputs with no extractable dosage, which GI treats as a separate category. Apart from this configuration, demographic differences are small relative to the framing effect.

Table 4: **Intra- and inter-vignette consistency of LLM dosage decisions using Gini impurity (GI) and KL divergence.** Global χ^2 is computed over the full race $\times$ gender contingency table and is non-significant at $\alpha=0.05$ for every configuration except Llama 3.1 (closed), which is significant at both tiers. Maximum KL divergence is at most 0.022 for every configuration except Llama 3.1 (closed) (0.079 at Tier 3). GI treats outputs with no extractable dosage as a separate category; KL divergence and χ^2 exclude them. GI values are not directly comparable across tiers because the number of options differs.

Model	Intra-vignette (Gini impurity)				Inter-vignette (KL)							
	Mean GI	% GI= 0	Mean GI	% GI= 0	Mean KL	Max KL	χ^2 p	V	Mean KL	Max KL	χ^2 p	V
	Tier 2 (binary)		Tier 3 (4-level)		Tier 2 (binary)				Tier 3 (4-level)			
GPT-4o-mini (closed)	0.037	87.5	0.006	98.0	0.0011	0.0046	0.838	0.047	0.0014	0.0063	0.922	0.048
GPT-4o-mini (RAG)	0.089	74.5	0.112	69.0	0.0017	0.0078	0.623	0.058	0.0059	0.0137	0.249	0.073
GPT-4.1-mini (closed)	0.047	86.0	0.118	68.0	0.0014	0.0040	0.750	0.052	0.0045	0.0104	0.521	0.064
GPT-4.1-mini (RAG)	0.000	100.0	0.028	92.0	0.0	0.0	1.0	NaN	0.0027	0.0061	0.309	0.072
GPT-5.4 (closed)	0.149	53.5	0.133	64.0	0.0013	0.0052	0.745	0.052	0.0023	0.0053	0.920	0.048
GPT-5.4 (RAG)	0.134	53.5	0.024	92.5	0.0047	0.0174	0.051	0.094	0.0056	0.0215	0.225	0.074
Llama 3.1 (closed)	0.227	48.0	0.404	13.0	0.0116	0.0591	0.004	0.144	0.0368	0.0790	< 0.001	0.157
Llama 3.1 (RAG)	0.087	78.5	0.035	90.5	0.0002	0.0005	0.999	0.021	0.0011	0.0028	0.997	0.034

4.3 Guideline RAG Results

RQ5: Guideline RAG substantially attenuates, but does not eliminate, the decision-space framing effect. Table 3 shows that retrieved guideline context first improves Tier-1 dosage accuracy under the original binary options and substantially raises accuracy for the two weakest closed-book configurations. On framing robustness, the mean accuracy drop is smaller under RAG (0.141 versus 0.212). Retention ratio also shows that every model retains a larger fraction of its accuracy under the expanded decision space when grounded, with the mean rising from 0.47 to 0.85, and absolute accuracy under framing is higher for all four. Mitigation is partial: GPT-4o-mini and Llama 3.1 still lose 0.397 and 0.255 accuracy points when the option set expands, against 0.038 and 0.05 gains for the other two.

Table 4 shows that the effect on demographic consistency is model-specific and does not follow the accuracy trend. Retrieval improves Tier-3 intra-vignette consistency for GPT-5.4, GPT-4.1-mini, and Llama-3.1, and reduces it for GPT-4o-mini. For Llama-3.1, retrieval also removes the significant demographic differences observed without it (Tier 3 χ^2 p = 0.997 with retrieval versus p < 0.001 without). Because guideline grounding (retrieval) also eliminates Llama-3.1’s None selections and nearly all of its outputs with no extractable dosage (Figures 1 and 5), part of the improvement in intra-vignette consistency may come from Llama-3.1 producing valid dosage answers more often. Retrieval therefore reduces the framing effect on average, while its effect on demographic consistency varies by model.

Models	TDP Acc.	Dosage Acc.	R-L	ES	LLMJ
Random Chunk Order					
GPT-4.1-mini	100.0	100.0	0.522	0.914	4.89
GPT-4o-mini	100.0	92.50	0.592	0.930	4.58
GPT-5.4	98.75	92.50	0.506	0.915	4.69
Llama-3.1	100.0	92.50	0.724	0.955	4.45
Reverse Chunk Order					
GPT-4.1-mini	100.0	100.0	0.557	0.920	4.91
GPT-4o-mini	100.0	95.00	0.577	0.933	4.68
GPT-5.4	86.25	92.5	0.426	0.792	4.19
Llama-3.1	97.50	95.0	0.630	0.923	4.33

Table 5: Performance comparison of Guideline RAG under random versus reverse chunk ordering. ‘TDP Acc.’ stands for Treatment Denial Prediction Accuracy (Yes/No prescribing decision); Dosage Acc.: dosage accuracy. Explanation quality is measured by R-L: Rouge-L (median), ES: embedding similarity, LLMJ: LLM-as-a-judge (out of 5).

RQ6: GPT-5.4 with RAG is sensitive to retrieved passage order. Guideline-augmented LLMs have a fixed set of retrieved passages prepended to each vignette (Table 1). Extending Q-PAIN’s counterfactual strategy to retrieved evidence, we re-present the same passages in reverse and random order; recommendations that reflect evidence content rather than presentation cues should be unaffected. For GPT-4.1-mini, GPT-4o-mini and Llama-3.1, treatment-denial prediction (TDP) accuracy stays at or above 97.5% under both permutations, and dosage accuracy ranges from 92.5% to 100% (Table 5). GPT-5.4’s TDP accuracy falls to 86.25% under reverse order (98.75% under random order), and its embedding similarity falls to 0.792, so its retrieval-conditioned reasoning is comparatively brittle to chunk-order perturbation.

5 Conclusion

Binary treatment-denial audits no longer distinguish current LLMs in post-operative pain management: all four models we tested approve treatment on nearly every Q-PAIN vignette. Their dosage decisions are much less stable. When we add a Medium (2 weeks) option and a None option to otherwise unchanged vignettes, dosage accuracy against the Q-PAIN expert label falls by 0.21 on average without retrieval and by 0.14 with guideline retrieval. Redefining Medium as 3 weeks reduces its uptake by 7.2 percentage points on average, so both its position and its duration influence selection. Demographic differences are small once clinical risk factors are added, except for closed-book Llama-3.1.

For anyone auditing a clinical LLM, the practical lesson is that a clean fairness result does not mean the model is stable. Demographic parity differences stayed at or below 0.10 in every GPT configuration with clinical risk factors, yet accuracy fell by up to 0.4 when only the answer options changed. We recommend that audits vary the answer format as well as patient identity, and score the resulting changes against ground truth wherever a label exists. Guideline retrieval helps, raising mean accuracy retention from 0.47 to 0.85, but it left GPT-4o-mini and Llama 3.1 with large losses, and GPT-5.4’s treatment decisions changed when the retrieved passages were presented in reverse order. Retrieval therefore needs its own robustness checks and cannot be assumed to fix the problem. Guideline grounding reduces the framing effect but does not remove it.

Future Work. First, collecting expert dosage labels for a sample of Tier 2 and Tier 3 risk-factor combinations would let us score the framing effect against ground truth beyond Tier 1. Second, randomizing the order in which dosage options are listed would separate attraction to the listed middle position from attraction to the intermediate duration itself. Third, we plan to replace flat chunk retrieval with structured guideline representations that encode scope, target population, and exclusion criteria, and to test whether this reduces sensitivity to passage order. We also plan to extend the framework to the four remaining Q-PAIN pain contexts.

Limitations

Our study has several limitations. First, we use only the ten post-operative vignettes of Q-PAIN, so

our findings may not generalize to its other pain contexts. Second, every Tier-1 vignette has the ground-truth label ‘Low’, so Tier-1 accuracy equals the rate of choosing ‘Low’. Tier 2 and Tier 3 results measure consistency rather than clinical correctness because expert dosage levels are unavailable for the risk-factor combinations; collecting these expert annotations forms our immediate future work. Third, dosage options are always listed in the same order and each prompt was run once, so we cannot separate the effect of the option’s position from that of its duration, or estimate run-to-run variability. Fourth, we assess the clinical relevance of retrieved passages only indirectly, through downstream dosage accuracy. Finally, Llama-3.1 (closed-book) produced no extractable dosage for 25.4% of Tier 3 items (11.9% in the sensitivity run), which limits comparison with the other configurations.

Ethical Considerations

We only extend the vignettes published by the original Q-PAIN authors (Logé et al., 2021), and these vignettes are hypothetical scenarios that do not include any personally identifiable information of a patient. The extended vignettes with added risk factors (mental health history, opioid tolerant status) were reviewed by our clinical collaborators, who are surgeons and trauma care physicians based in the US, to ensure that they do not perpetuate stigmatizing associations. We acknowledge the risk that individual metrics, such as a low cross-vignette KL Divergence score, used in isolation may lead to misleading inferences. In our experiments, demographic parity differences are small for the GPT models once clinical risk factors are added, while accuracy under the expanded option set drops substantially. We therefore report label, explanation, and retrieval metrics together, and recommend that audits using our framework do the same.

Acknowledgments

This work was supported by the Agency for Health Research & Quality under Award Number R01HS027434. The content is solely the responsibility of the authors and does not necessarily represent the official views of the funders.

References

Yaara Artsi, Vera Sorin, Benjamin S Glicksberg, Panagiotis Korfiatis, Girish N Nadkarni, and Eyal Klang.

2025. Large language models in real-world clinical workflows: a systematic review of applications and implementation. *Front. Digit. Health*, 7(1659134):1659134.
- Gunjan Balde, Soumyadeep Roy, Mainack Mondal, and Niloy Ganguly. 2025. [Evaluation of LLMs in medical text summarization: The role of vocabulary adaptation in high OOV settings](#). In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 22989–23004, Vienna, Austria. Association for Computational Linguistics.
- Gunjan Balde, Soumyadeep Roy, Mainack Mondal, and Niloy Ganguly. 2026. [Learning faster with better tokens: Parameter-efficient vocabulary adaptation for specialized text summarization](#). In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 45071–45086, San Diego, California, United States. Association for Computational Linguistics.
- Helene Beloeil, Girish P Joshi, Patricia Lavand’homme, Felicia Cox, Jennifer Summers, Esther Pogatzki-Zahn, and Patrice Forget. 2025. Perioperative analgesia in high-risk patients, what clinicians and searchers must know? a narrative review. *Anaesth Crit Care Pain Med*, 45(3):101644.
- Hanjie Chen, Zhouxiang Fang, Yash Singla, and Mark Dredze. 2025. [Benchmarking large language models on answering and explaining challenging medical questions](#). In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 3563–3599, Albuquerque, New Mexico. Association for Computational Linguistics.
- Emma Croxford, Yanjun Gao, Elliot First, Nicholas Pellegrino, Miranda Schnier, John Caskey, Madeline Oguss, Graham Wills, Guanhua Chen, Dmitriy Dligach, Matthew M Churpek, Anoop Mayampurath, Frank Liao, Cherodeep Goswami, Karen K. Wong, Brian W. Patterson, and Majid Afshar. 2025. [Automating evaluation of AI text generation in healthcare with a large language model \(LLM\)-as-a-judge](#). *medRxiv*.
- Deborah Dowell, Kathleen R Ragan, Christopher M Jones, Grant T Baldwin, and Roger Chou. 2022. CDC clinical practice guideline for prescribing opioids for pain - united states, 2022. *MMWR Recomm. Rep.*, 71(3):1–95.
- Yanjun Gao, Dmitriy Dligach, Timothy Miller, and Majid Afshar. 2023. [Overview of the problem list summarization \(ProbSum\) 2023 shared task on summarizing patients’ active diagnoses and problems from electronic health record progress notes](#). In *Proceedings of the 22nd Workshop on Biomedical Natural Language Processing and BioNLP Shared Tasks*, pages 461–467, Toronto, Canada. Association for Computational Linguistics.
- Akhil Gourabathina, Walter Gerych, Evan Pan, and Marzyeh Ghassemi. 2025. [The medium is the message: How non-clinical information shapes clinical decisions in LLMs](#). In *Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency*, pages 1805–1828. ACM.
- Kelly M. Hoffman, Sophie Trawalter, Jordan R. Axt, and M. Norman Oliver. 2016. Racial bias in pain assessment and treatment recommendations, and false beliefs about biological differences between blacks and whites. *Proceedings of the National Academy of Sciences*, 113(16):4296–4301.
- Joel Huber, John W. Payne, and Christopher Puto. 1982. [Adding asymmetrically dominated alternatives: Violations of regularity and the similarity hypothesis](#). *Journal of Consumer Research*, 9(1):90–98.
- Paulyne Lee, Maxine Le Saux, Rebecca Siegel, Monika Goyal, Chen Chen, Yan Ma, and Andrew C. Meltzer. 2019. Racial and ethnic disparities in the management of acute pain in US emergency departments: Meta-analysis and systematic review. *The American Journal of Emergency Medicine*, 37(9):1770–1777.
- Cécile Logé, Emily Ross, David Yaw Amoah Dadey, Saahil Jain, Adriel Saporta, Andrew Y. Ng, and Pranav Rajpurkar. 2021. [Q-Pain: A question answering dataset to measure social bias in pain management](#). In *Proceedings of the 35th Conference on Neural Information Processing Systems (NeurIPS 2021), Track on Datasets and Benchmarks*.
- Sanjay Mohanty, Ronnie A Rosenthal, Marcia M Russell, Mark D Neuman, Clifford Y Ko, and Nestor F Esnaola. 2016. Optimal perioperative management of the geriatric patient: A best practices guideline from the american college of surgeons NSQIP and the american geriatrics society. *J. Am. Coll. Surg.*, 222(5):930–947.
- Mahmud Omar, Shelly Soffer, Reem Agbareia, Nicola Luigi Bragazzi, Benjamin S. Glicksberg, Yasmin L. Hurd, Donald U. Apakama, Alexander W. Charney, David L. Reich, Girish N. Nadkarni, and Eyal Klang. 2026. [Socio-demographic gaps in pain management guided by large language models](#). *Nature Health*, 1(2):216–225.
- Jesutofunmi A. Omiye, Jenna C. Lester, Simon Spichak, Veronica Rotemberg, and Roxana Daneshjou. 2023. [Large language models propagate race-based medicine](#). *npj Digital Medicine*, 6(1):195.
- Pouya Pezeshkpour and Estevam Hruschka. 2024. [Large language models sensitivity to the order of options in multiple-choice questions](#). In *Findings of the Association for Computational Linguistics: NAACL 2024*, pages 2006–2017, Mexico City, Mexico. Association for Computational Linguistics.
- Anke Samulowitz, Ida Gremyr, Erik Eriksson, and Gunnel Hensing. 2018. “brave men” and “emotional women”: A theory-guided literature review on gender bias in health care and gendered norms towards

- patients with chronic pain. *Pain Research and Management*, 2018:1–14.
- Krithi Shailya, Shreya Rajpal, Gokul S Krishnan, and Balaraman Ravindran. 2025. [Lext: Towards evaluating trustworthiness of natural language explanations](#). In *Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency, FAccT '25*, page 1565–1587, New York, NY, USA. Association for Computing Machinery.
- Itamar Simonson. 1989. [Choice based on reasons: The case of attraction and compromise effects](#). *Journal of Consumer Research*, 16(2):158–174.
- Karan Singhal, Shekoofeh Azizi, Tao Tu, S. Sara Mahdavi, Jason Wei, Hyung Won Chung, Nathan Scales, Ajay Tanwani, Heather Cole-Lewis, Stephen Pfohl, and 1 others. 2023. Large language models encode clinical knowledge. *Nature*, 620(7972):172–180.
- Annalisa Szymanski, Noah Ziemis, Heather A. Eicher-Miller, Toby Jia-Jun Li, Meng Jiang, and Ronald A. Metoyer. 2025. [Limitations of the llm-as-a-judge approach for evaluating llm outputs in expert knowledge tasks](#). In *Proceedings of the 30th International Conference on Intelligent User Interfaces, IUI '25*, page 952–966, New York, NY, USA. Association for Computing Machinery.
- Konstantinos Tzioumis. 2018. Demographic aspects of first names. *Scientific Data*, 5(1):1–9.
- Carol S. Weisse, Paul C. Sorum, Kafi N. Sanders, and Beth L. Syat. 2001. Do gender and race affect decisions about pain management? *Journal of General Internal Medicine*, 16(4):211–217.
- Timothy H. Wideman, Robert R. Edwards, David M. Walton, Marc O. Martel, Anne Hudon, and David A. Seminowicz. 2019. The multimodal assessment model of pain: A novel framework for further integrating the subjective pain experience within research and practice. *The Clinical Journal of Pain*, 35(3):212.
- Cameron C Young, Elizabeth Enichen, Arya Rao, and Marc D Succ. 2025. Racial, ethnic, and sex bias in large language model opioid recommendations for pain management. *Pain*, 166(3):511–517.
- Travis Zack, Eric Lehman, Mirac Suzgun, Jorge A. Rodriguez, Leo Anthony Celi, Judy Gichoya, Dan Jurafsky, Peter Szolovits, David W. Bates, Raja-Elie E. Abdunour, and 1 others. 2024. Assessing the potential of GPT-4 to perpetuate racial and gender biases in health care: A model evaluation study. *The Lancet Digital Health*, 6(1):e12–e22.
- Haoran Zhang, Amy X. Lu, Mohamed Abdalla, Matthew McDermott, and Marzyeh Ghassemi. 2020. Hurtful words: Quantifying biases in clinical contextual word embeddings. In *Proceedings of the ACM Conference on Health, Inference, and Learning (CHIL)*, pages 110–120. ACM.
- Chujie Zheng, Hao Zhou, Fandong Meng, Jie Zhou, and Minlie Huang. 2024. [Large language models are not robust multiple choice selectors](#). In *The Twelfth International Conference on Learning Representations (ICLR)*.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, Hao Zhang, Joseph Gonzalez, and Ion Stoica. 2023. [Judging llm-as-a-judge with mt-bench and chatbot arena](#). In *Advances in Neural Information Processing Systems*, volume 36, pages 46595–46623. Curran Associates, Inc.

A Implementation Details

The closed-source OpenAI models include GPT-5.4 (a reasoning model, with a model version date of 2026-03-05), GPT-4.1-mini (model version date 2025-04-14), and GPT-4o-mini (model version date: 2024-07-18), while the open-weight model is Llama 3.1 (8B)², run locally using Ollama. We aim to include models that differ in recency and accessibility. The parameter counts of the OpenAI models are not publicly disclosed; Llama 3.1 has 8 billion parameters. Given the fast pace of LLM model updates, we include GPT-5.4 (2026).

Computational Resources. All open-weight model inference was conducted locally on a MacBook Pro equipped with an Apple M4 Pro chip (14-core CPU, 20-core GPU, 48 GB unified memory), without access to external GPU clusters. The open-weight model, Llama 3.1-8B, is served via Ollama v0.17.4. The same Ollama instance hosted the qwen3-embedding model (8B variant) for two purposes: (i) dense passage retrieval in the guideline RAG pipeline, and (ii) similarity scoring between ground-truth and model-generated clinical rationales. The LLM judge, MedGemma-27B, was likewise served locally via Ollama. Proprietary models (GPT-4o mini, GPT-4.1 mini, GPT-5.4) were queried through the OpenAI API. All local inference was executed sequentially on the host machine; no cloud or high-performance computing infrastructure was used for open-weight model evaluation. Table 2 describes the computational budget in terms of the number of inference calls.

B Demographic Parity Difference

We evaluate every model setting using the standard fairness metric implemented in the fairlearn

²<https://ollama.com/library/llama3.1:8b>

Python package³: *Demographic Parity Difference* (DPD; $\max_g \text{SelectionRate}(\text{SR})_g - \min_g \text{SR}_g$). DPD is a parity metric and does not use ground truth. DPD is computed on the ‘predict-Low’ outcome, the most directly comparable binary target across tiers (it is the only valid binary outcome at Tier 1, and the principal dosage class whose distribution shifts at Tier 3). Because every Tier 1 vignette has ground truth Low, equal-opportunity and equalized-odds differences reduce to DPD at Tier 1, so we report DPD only.

Findings. Figure 4 reports DPD on the predict-Low outcome across Tiers 1–3 for the six GPT configurations. At Tiers 2 and 3, intersectional DPD is at most 0.10 for every configuration (maximum 0.100, GPT-4.1-mini closed-book at Tier 3), consistent with the non-significant χ^2 tests in Table 4. Tier 1 values are larger for some configurations (0.40 for GPT-5.4 closed-book, 0.20 for GPT-4o-mini closed-book), but each intersectional group contains only 10 decisions at Tier 1, so these values move in steps of 0.10. Comparing Tier 2 with Tier 3, which differ only in the dosage options, GPT-4.1-mini’s intersectional DPD rises to 0.10, whereas GPT-5.4 closed-book’s changes little; GPT-5.4’s larger drop occurs between Tiers 1 and 2, when clinical risk factors are added. DPD can therefore change with the option set even when patient demographics do not.

C Experimental Results

C.1 Demographic Differences in Tier 2 to Tier 3 Escalation and De-escalation (GPT-4o-mini)

Figure 6 shows that when the Medium option becomes available, GPT-4o-mini moves nearly all of its Tier 2 High assignments to a lower dosage and escalates a subset of its few Tier 2 Low assignments, with the only demographic effect reaching significance being a small gender gap in de-escalation. Consistent with Young et al. (2025) and Omar et al. (2026), demographic differences in dosage are small.

D Guideline-RAG Corpus and Pipeline

Our evidence corpus comprises 77 perioperative pain management guidelines: (i) 11 perioperative care guidelines recommended by clinicians including CDC 2022 (Dowell et al., 2022), ACS

NSQIP/AGS (Mohanty et al., 2016), NICE NG180⁴, and Surgery and Opioids 2021, processed with GROBID into XML and then converted to structured text, and (ii) PubMed ‘pain management’ and ‘perioperative care’ guideline articles available from the PMC Open Access subset (last 10 years). You can reproduce the search results accessed on December 23, 2025, by the following url: https://pubmed.ncbi.nlm.nih.gov/?term=pubmed+pmc+open+access%5Bfilter%5D+%22pain+management%22&filter=datasearch.y_10&filter=pubt.practiceguideline&size=50. We downloaded the articles in .xml format and parsed them. We removed the following tags and their content: table-wrap, fig, graphic, media, supplementary-material, ref-list, ref, xref, fn, fn-group. We also removed sections whose titles included: appendix, supplementary, additional file, electronic supplementary, acknowledg, funding, author contributions, competing interests, references, bibliography. After the above cleaning step, the textual content is available in both XML and plaintext format.

E Generative AI Usage Statement

We used generative AI assistants to support writing tasks, specifically for paraphrasing and refining text. All outputs were carefully reviewed and validated by the authors. A coding assistant (Copilot) was used to write plotting code.

³<https://pypi.org/project/fairlearn/>

⁴<https://www.nice.org.uk/guidance/ng180/chapter/Recommendations>




Figure 4: Demographic Parity Difference on the predict-Low outcome across Tiers 1, 2, and 3 for the six GPT settings with complete factorial coverage, faceted by sensitive-attribute granularity. Tier 2 and Tier 3 intersectional DPDs are at or below 0.10, consistent with Table 4. Tier 1 values rest on 10 decisions per intersectional group.

System Prompt

You are a board-certified anaesthesiologist and clinical-reasoning expert. Your task is to evaluate how well a candidate clinical explanation captures the same medical content as a reference explanation for a post-operative pain management vignette.

Score the candidate on a scale of 1–5 using ONLY these criteria:

- 5 – Excellent: all key clinical facts, reasoning, and recommendations match.
- 4 – Good: most key points match; minor omissions or wording differences only.
- 3 – Adequate: core reasoning is present but notable gaps or inaccuracies exist.
- 2 – Poor: partial overlap; key clinical facts are missing or incorrect.
- 1 – Unacceptable: the candidate fails to capture the clinical content.

Respond with a SINGLE integer (1, 2, 3, 4, or 5) and nothing else.

User Prompt (per vignette, with placeholders filled during runtime)

Reference explanation:

<Replace with ground-truth explanation>

Candidate explanation:

<Replace with candidate explanation>

Score (1 - 5):

Table 6: LLM-as-a-judge prompt used with MedGemma-27B to grade the similarity between ground-truth explanation and LLM-generated explanation for Tier-1 Q-PAIN experiments. The system message and user message were passed as separate turns in the Ollama /api/chat interface.




Figure 5: Tier 3 dosage decisions (four-level options, Medium = 2 weeks) for the Tier 2 items. No ground-truth labels are available because the added risk factors may change the appropriate dosage. Unknown: no dosage could be extracted from the output.




Figure 6: GPT-4o-mini Tier 2 to Tier 3 escalation and de-escalation rates by gender and race (95% Wilson intervals). Escalation rates are noisy due to sparse Tier 2 Low prescriptions. De-escalation rates are near-saturated.